

# Grashof\_David\_Assignment1

April 7, 2019

## 0.1 Assignment 1 | Exploring and Visualizing Data

### 0.1.1 Overview

The Northwestern MSDS program administrators are interested in finding out the best way to structure the program, going

forward. More specifically, they would like to know what programming language to teach the courses in and what electives

they will offer. To make a fully-informed decision, a survey was distributed to current students in the MSDS program to

gauge their interests in both programming languages and potential electives. The survey has been collected and analyzed in

various ways to help guide the administrators in enhancing the program.

In order to address each question, the data was parsed out accordingly from the original file.

To answer the first

question, which programming language should be used primarily for future classes, a mean was taken for all responses. Due

to high level of variance, this value was not very meaningful. What it did show was that both R and Python had the highest

average response in the mid-30s. To get a better understanding of the responses, a histogram was constructed for each

language. It became increasingly obvious that both Java languages and SAS could be removed from consideration, leaving only

R and Python. Before making a final determination, it was necessary to understand if the responses changed based on the

context, such as industry, professional or personal. A correlation plot showed that the responses between each context for

a given language were highly correlated between .65 and .79, indicating that students mostly had the same response

regardless of context. Another consideration was the determine if there was a link between courses completed and responses,

with the assumption being that the answers of students with more classes completed would carry more weight than new

students who may not have any programming experience. The results showed that there was no relationship.

To determine which classes should be offered, first the mean was calculated for the four possible options with the Python

course logging the highest mean score of 73.5, much greater than the second highest class, Foundations in Data Engineering, with a mean score of 58. Likewise, this data exhibited high variance like the language preference responses. Next, the distributions of the responses for each course were plotted, further showing that the Python class had the most responses at or near 100. It is recommended that given the high interest shown for Python in both questions, that the Python class be included as an elective. Additionally, given the similar scores shown for both R and Python as the preferred programming language, future courses should be offered as one or the other.

### 0.1.2 1.1 | Load Modules

```
In [1]: import pandas as pd # data frame operations
import nbconvert # convert to pdf
import numpy as np # arrays and math functions
import matplotlib.pyplot as plt # static plotting
import seaborn as sns # pretty plotting, including heat map
```

### 0.1.3 1.2 | Load Data

```
In [2]: #define filepath
file = "C:/Users/David/OneDrive/MSDS/MSDS422/Week1/Assignment/mspa-software-survey-cas
df = pd.read_csv(file)
#set RespondentID as index
df.set_index('RespondentID', drop = True, inplace = True)
#check structure of data
df.head()
```

```
Out [2]:
```

	Personal_JavaScalaSpark	Personal_JavaScriptHTMLCSS \
RespondentID		
5135740122	0	0
5133300037	10	10
5132253300	20	0
5132096630	10	10
5131990362	20	0

	Personal_Python	Personal_R	Personal_SAS \
RespondentID			
5135740122	0	50	50
5133300037	50	30	0
5132253300	40	40	0
5132096630	25	35	20
5131990362	0	70	10

	Professional_JavaScalaSpark	Professional_JavaScriptHTMLCSS \
--	-----------------------------	----------------------------------

RespondentID

5135740122	0	0
5133300037	25	25
5132253300	0	0
5132096630	10	10
5131990362	20	0

	Professional_Python	Professional_R	Professional_SAS	...	\
RespondentID				...	
5135740122	0	25	75	...	
5133300037	30	20	0	...	
5132253300	40	40	20	...	
5132096630	25	35	20	...	
5131990362	0	80	0	...	

	PREDICT453	PREDICT454	PREDICT455	PREDICT456	PREDICT457	\
RespondentID						
5135740122	NaN	NaN	NaN	NaN	NaN	
5133300037	NaN	NaN	NaN	NaN	NaN	
5132253300	NaN	NaN	NaN	NaN	NaN	
5132096630	NaN	NaN	NaN	NaN	NaN	
5131990362	NaN	NaN	NaN	NaN	NaN	

	OtherPython	OtherR	OtherSAS	Other	Graduate_Date
RespondentID					
5135740122	NaN	NaN	NaN	NaN	NaN
5133300037	NaN	NaN	NaN	NaN	Spring 2018
5132253300	NaN	NaN	NaN	NaN	Fall 2018
5132096630	NaN	NaN	NaN	NaN	Fall 2017
5131990362	NaN	NaN	NaN	CS-435 with Weka	Fall 2018

[5 rows x 40 columns]

### 0.1.4 1.3 | Pre-processing Data

In [3]: *#organize fields into easily callable groups by purpose of language*

```
personal_fields = ["Personal_JavaScalaSpark", "Personal_JavaScriptHTMLCSS", "Personal_Py
```

```
professional_fields = ["Professional_JavaScalaSpark", "Professional_JavaScriptHTMLCSS",
```

```
industry_fields = ["Industry_JavaScalaSpark", "Industry_JavaScriptHTMLCSS", "Industry_Py
```

*#organize fields into easily callable by language regardless of purpose of use*

```
jss_fields = ["Personal_JavaScalaSpark", "Professional_JavaScalaSpark", "Industry_JavaSc
```

```
jsh_fields = ["Personal_JavaScriptHTMLCSS", "Professional_JavaScriptHTMLCSS", "Industry_
```

```
python_fields = ["Personal_Python", "Professional_Python", "Industry_Python"]
```

```
r_fields = ["Personal_R", "Professional_R", "Industry_R"]
```

```
sas_fields = ["Personal_SAS", "Professional_SAS", "Industry_SAS"]
```

*#organize classes taken fields*

```
class_fields = ["PREDICT400", "PREDICT401", "PREDICT410", "PREDICT411", "PREDICT413", "PRED
```

```
"PREDICT452", "PREDICT453", "PREDICT454", "PREDICT455", "PREDICT456", "PRED
```

```

#organize interest fields
interest_fields = ["Python_Course_Interest","Foundations_DE_Course_Interest","Analytic

# def encoder to create a binary nan value
def value_encoder(i):
    if i != i: catstr = "0"
    else: catstr = "1"
    catstr = int(catstr)
    return(catstr)

# apply encoder to class values in order to sum how many classes an individual has tak
for i in df.columns[20:38]:
    df[i] = df[i].map(value_encoder)

#sum number of programming classes taken per respondent as Courses_Completed may inclu
df["Courses_Completed_Prog"] = df[class_fields].sum(axis=1)

def courses_completed_group(i):
    i = int(i)
    if i > 0 and i <= 3: catstr = "0-3"
    elif i > 3 and i <= 6: catstr = "4-6"
    else: catstr = "7-10"
    return(catstr)

#map bucket
df["Courses_Completed_Bucket"] = df["Courses_Completed_Prog"].map(courses_completed_gr
#rebuild class organization of fields
class_fields = ["PREDICT400","PREDICT401","PREDICT410","PREDICT411","PREDICT413","PRED
                "PREDICT452","PREDICT453","PREDICT454","PREDICT455","PREDICT456","PRED

```

```

In [4]: #overview of key statistics
for i in [personal_fields,professional_fields,industry_fields]:
    print(df[i].describe())
    print('-----')

#on a personal level R and Python are the leaders by mean but all language have a high
#on a professional level R and Python are again the leaders by mean but with equally h
#on an industry level R and Pytho are again again the leaders by mean with high std

```

	Personal_JavaScalaSpark	Personal_JavaScriptHTMLCSS	Personal_Python \
count	207.000000	207.000000	207.000000
mean	10.135266	4.797101	31.304348
std	11.383477	6.757764	15.570982
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	20.000000
50%	9.000000	0.000000	30.000000
75%	20.000000	10.000000	40.000000
max	70.000000	30.000000	90.000000

	Personal_R	Personal_SAS
count	207.000000	207.000000
mean	37.125604	16.637681
std	14.576003	13.626400
min	0.000000	0.000000
25%	30.000000	5.000000
50%	35.000000	15.000000
75%	50.000000	25.000000
max	100.000000	75.000000

---

	Professional_JavaScalaSpark	Professional_JavaScriptHTMLCSS \
count	207.000000	207.000000
mean	9.251208	5.840580
std	13.167505	10.812555
min	0.000000	0.000000
25%	0.000000	0.000000
50%	5.000000	0.000000
75%	15.000000	10.000000
max	80.000000	100.000000

	Professional_Python	Professional_R	Professional_SAS
count	207.000000	207.000000	207.000000
mean	30.028986	36.415459	18.463768
std	19.144802	20.847606	18.831841
min	0.000000	0.000000	0.000000
25%	20.000000	25.000000	0.000000
50%	30.000000	33.000000	15.000000
75%	40.000000	50.000000	30.000000
max	100.000000	100.000000	100.000000

---

	Industry_JavaScalaSpark	Industry_JavaScriptHTMLCSS	Industry_Python \
count	207.000000	207.000000	207.000000
mean	11.942029	6.966184	29.772947
std	14.706399	10.030721	17.959816
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	20.000000
50%	5.000000	0.000000	30.000000
75%	20.000000	10.000000	40.000000
max	70.000000	50.000000	95.000000

	Industry_R	Industry_SAS
count	207.000000	207.000000
mean	32.434783	18.884058
std	15.912209	19.137623
min	0.000000	0.000000
25%	22.500000	0.000000
50%	30.000000	15.000000
75%	40.000000	30.000000

```
max      85.000000    100.000000
```

---

```
In [5]: for i in [jss_fields,jsh_fields,python_fields,r_fields,sas_fields]:
```

```
        print(df[i].describe())
```

```
        print('-----')
```

```
        # All 5 languages have near equal distribution of scores regardless of intent (personal vs professional)
```

```
        # disregarding intent in further analysis may be acceptable. Scaling these scores may vary by language
```

	Personal_JavaScalaSpark	Professional_JavaScalaSpark \
count	207.000000	207.000000
mean	10.135266	9.251208
std	11.383477	13.167505
min	0.000000	0.000000
25%	0.000000	0.000000
50%	9.000000	5.000000
75%	20.000000	15.000000
max	70.000000	80.000000

	Industry_JavaScalaSpark
count	207.000000
mean	11.942029
std	14.706399
min	0.000000
25%	0.000000
50%	5.000000
75%	20.000000
max	70.000000

---

	Personal_JavaScriptHTMLCSS	Professional_JavaScriptHTMLCSS \
count	207.000000	207.000000
mean	4.797101	5.840580
std	6.757764	10.812555
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	10.000000	10.000000
max	30.000000	100.000000

	Industry_JavaScriptHTMLCSS
count	207.000000
mean	6.966184
std	10.030721
min	0.000000
25%	0.000000
50%	0.000000
75%	10.000000

max 50.000000

---

	Personal_Python	Professional_Python	Industry_Python
count	207.000000	207.000000	207.000000
mean	31.304348	30.028986	29.772947
std	15.570982	19.144802	17.959816
min	0.000000	0.000000	0.000000
25%	20.000000	20.000000	20.000000
50%	30.000000	30.000000	30.000000
75%	40.000000	40.000000	40.000000
max	90.000000	100.000000	95.000000

---

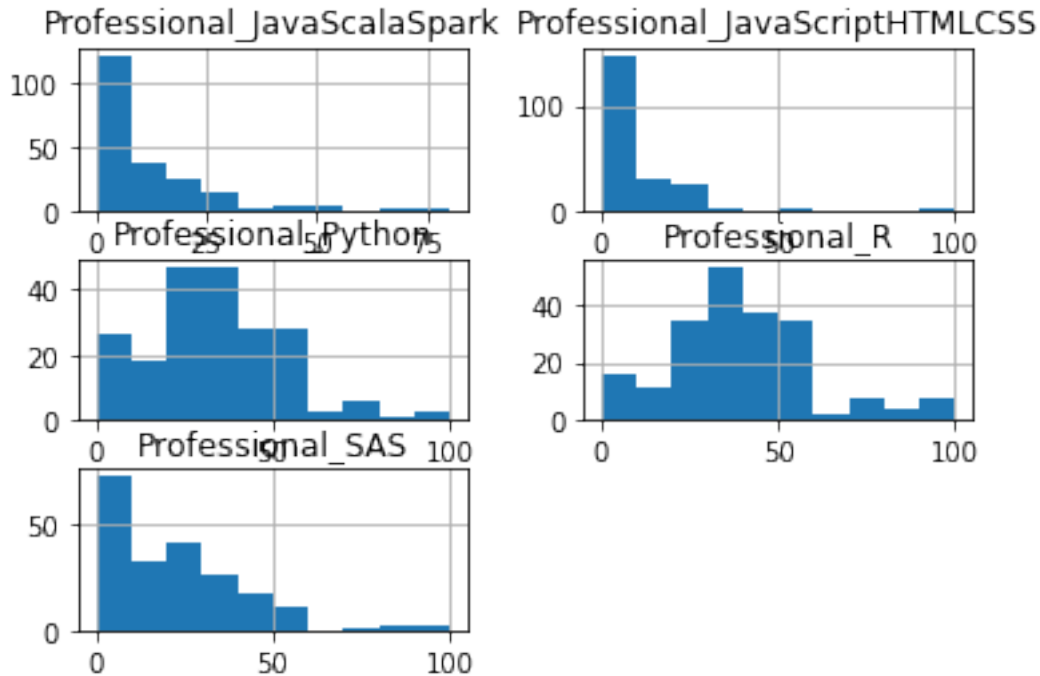
	Personal_R	Professional_R	Industry_R
count	207.000000	207.000000	207.000000
mean	37.125604	36.415459	32.434783
std	14.576003	20.847606	15.912209
min	0.000000	0.000000	0.000000
25%	30.000000	25.000000	22.500000
50%	35.000000	33.000000	30.000000
75%	50.000000	50.000000	40.000000
max	100.000000	100.000000	85.000000

---

	Personal_SAS	Professional_SAS	Industry_SAS
count	207.000000	207.000000	207.000000
mean	16.637681	18.463768	18.884058
std	13.626400	18.831841	19.137623
min	0.000000	0.000000	0.000000
25%	5.000000	0.000000	0.000000
50%	15.000000	15.000000	15.000000
75%	25.000000	30.000000	30.000000
max	75.000000	100.000000	100.000000

---

```
In [6]: #plot histogram for all defined professional fields
df[professional_fields].hist()
plt.xlabel("Interest Score (0-100)")
plt.grid(True)
plt.show()
```



In [7]: *#visualizing scatterplots to confirm similiarities of scores across intent*

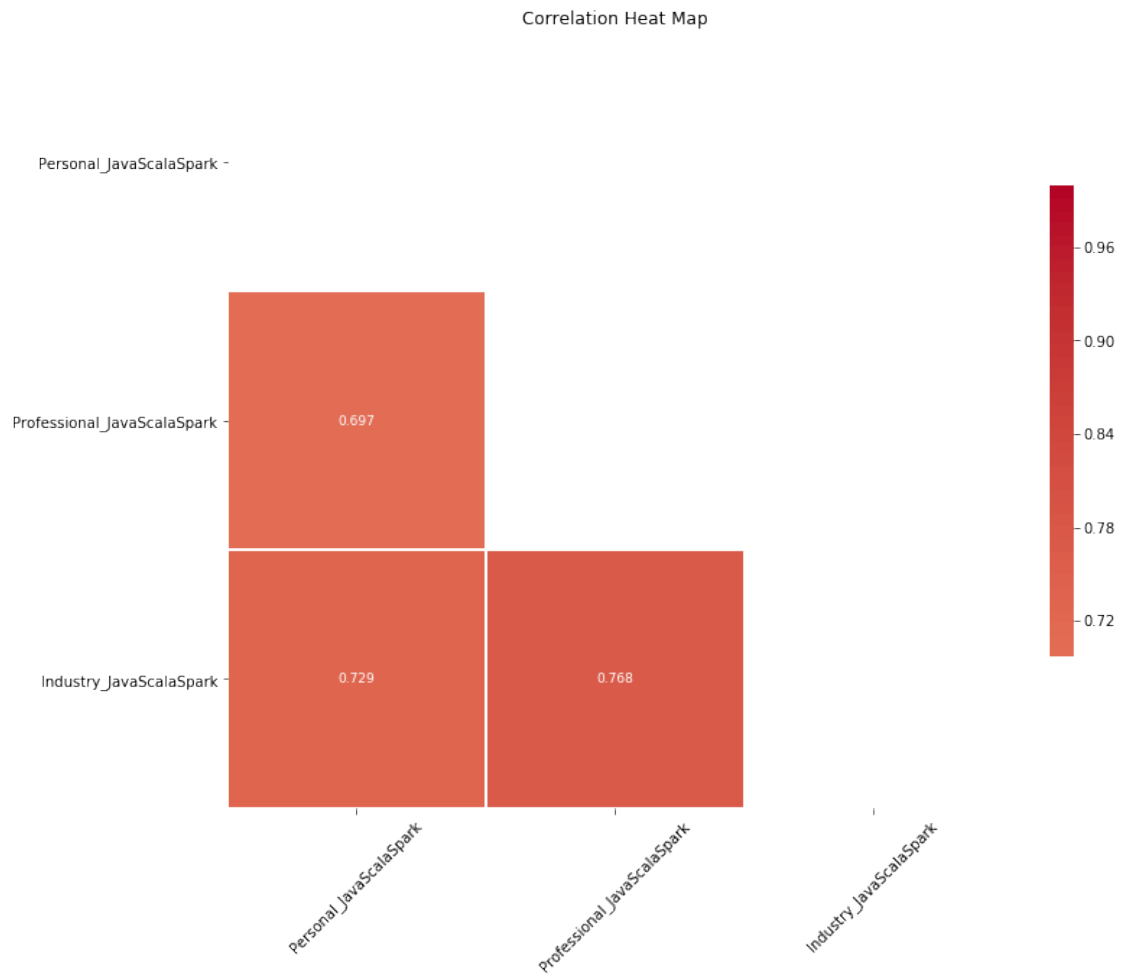
```
def corr_chart(df_corr):
    corr=df_corr.corr()
    #screen top half to get a triangle
    top = np.zeros_like(corr, dtype=np.bool)
    top[np.triu_indices_from(top)] = True
    fig=plt.figure()
    fig, ax = plt.subplots(figsize=(12,12))
    sns.heatmap(corr, mask=top, cmap='coolwarm',
                center = 0, square=True,
                linewidths=.5, cbar_kws={'shrink':.5},
                annot = True, annot_kws={'size': 9}, fmt = '.3f')
    plt.xticks(rotation=45) # rotate variable labels on columns (x axis)
    plt.yticks(rotation=0) # use horizontal variable labels on rows (y axis)
    plt.title('Correlation Heat Map')
    plt.savefig('plot-corr-map.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)

for i in [jss_fields,jsh_fields,python_fields,r_fields,sas_fields]:
    corr_chart(df_corr = df[i])
```

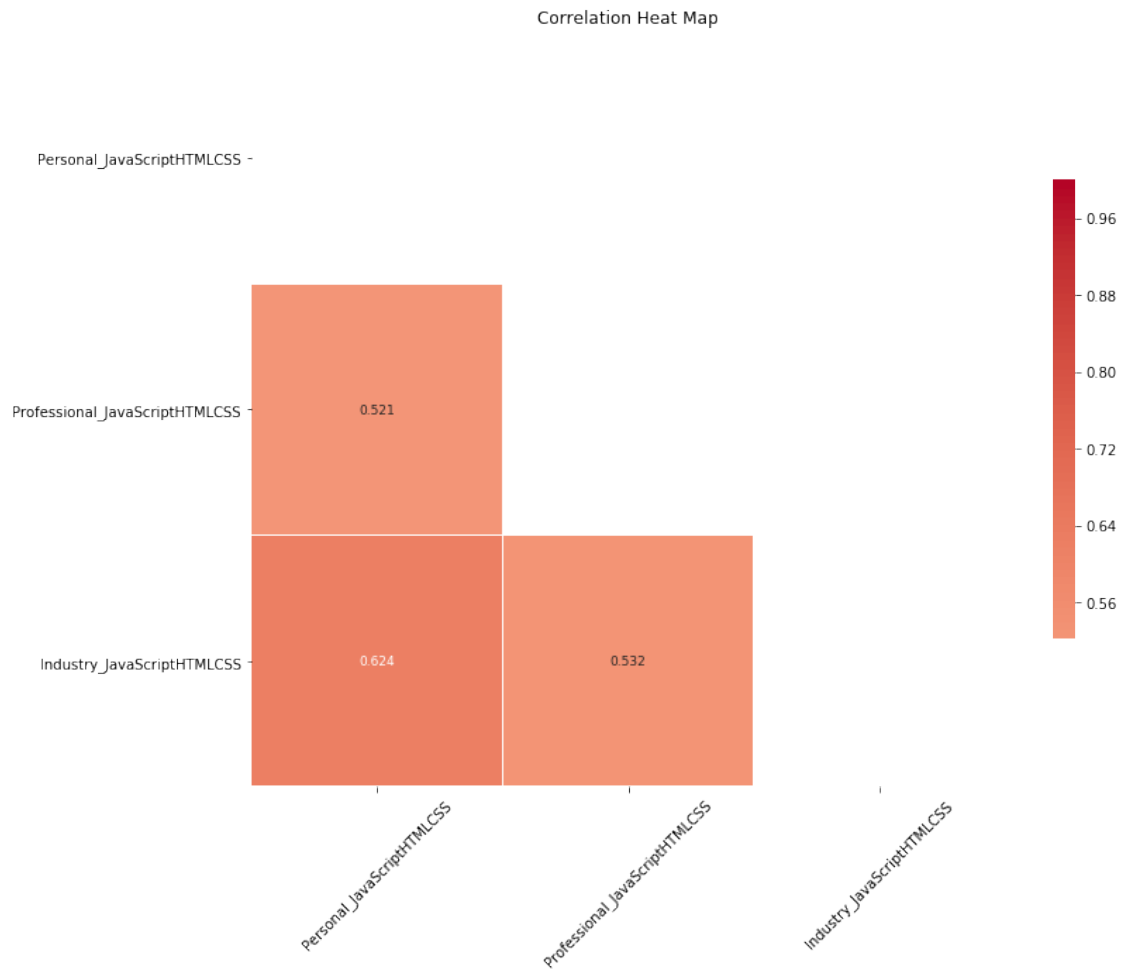
*#chart appears to indicate that language by intent is not as highly correlated as prev*  
*#JavascriptHTML has the weakest correlation by intent at .52 - .62 while Python shows*



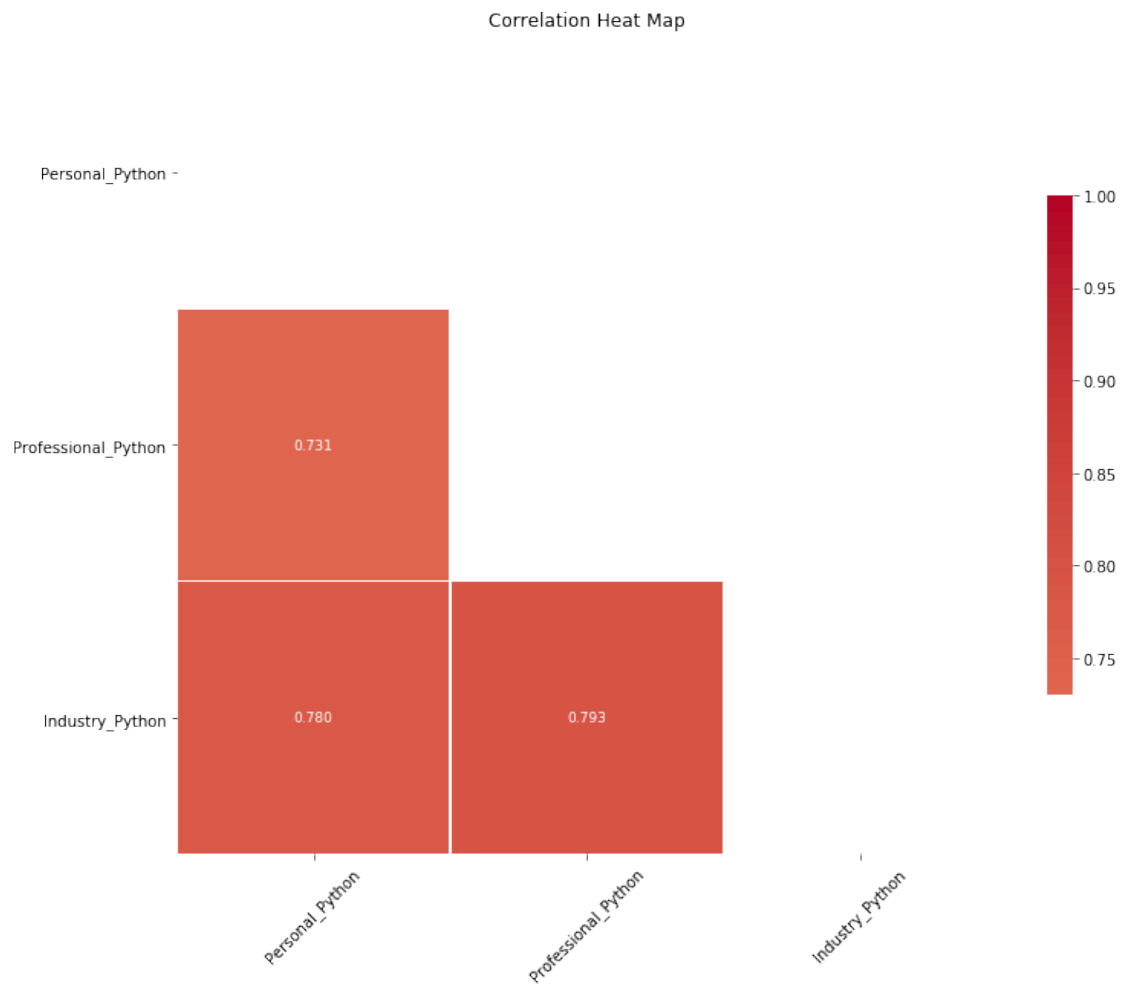
<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>

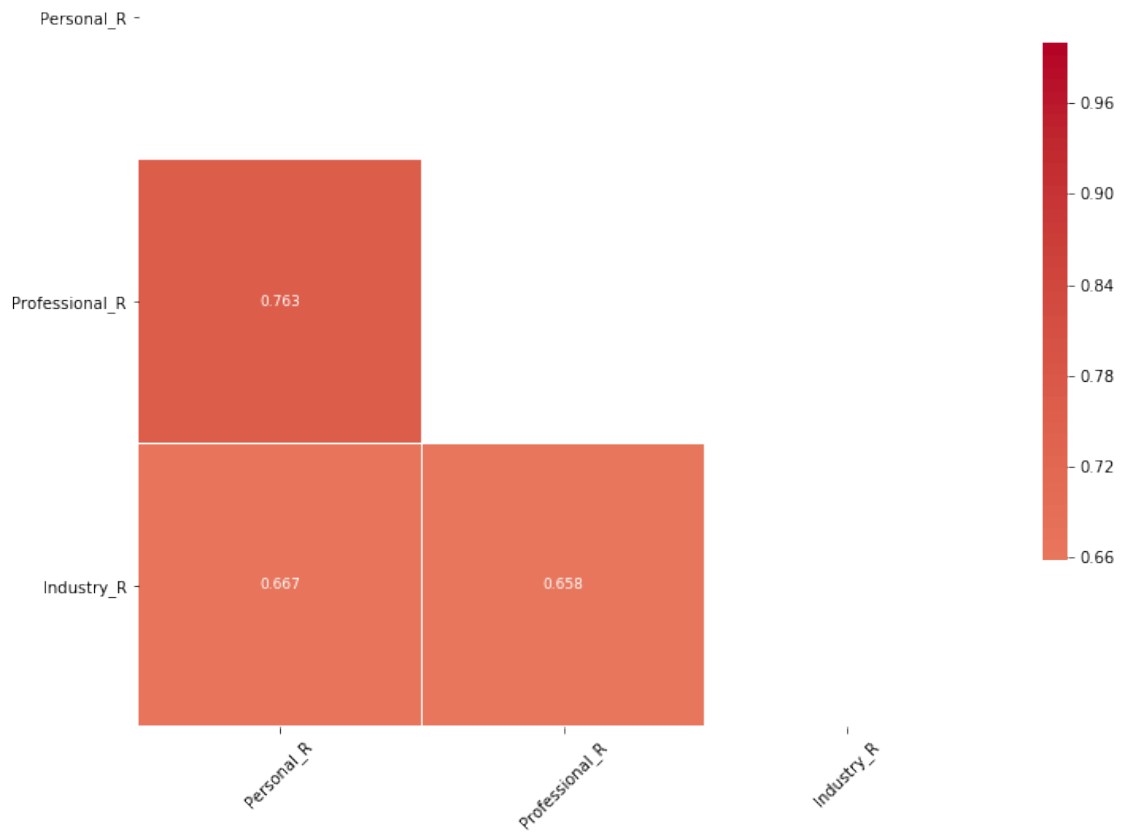


<Figure size 432x288 with 0 Axes>

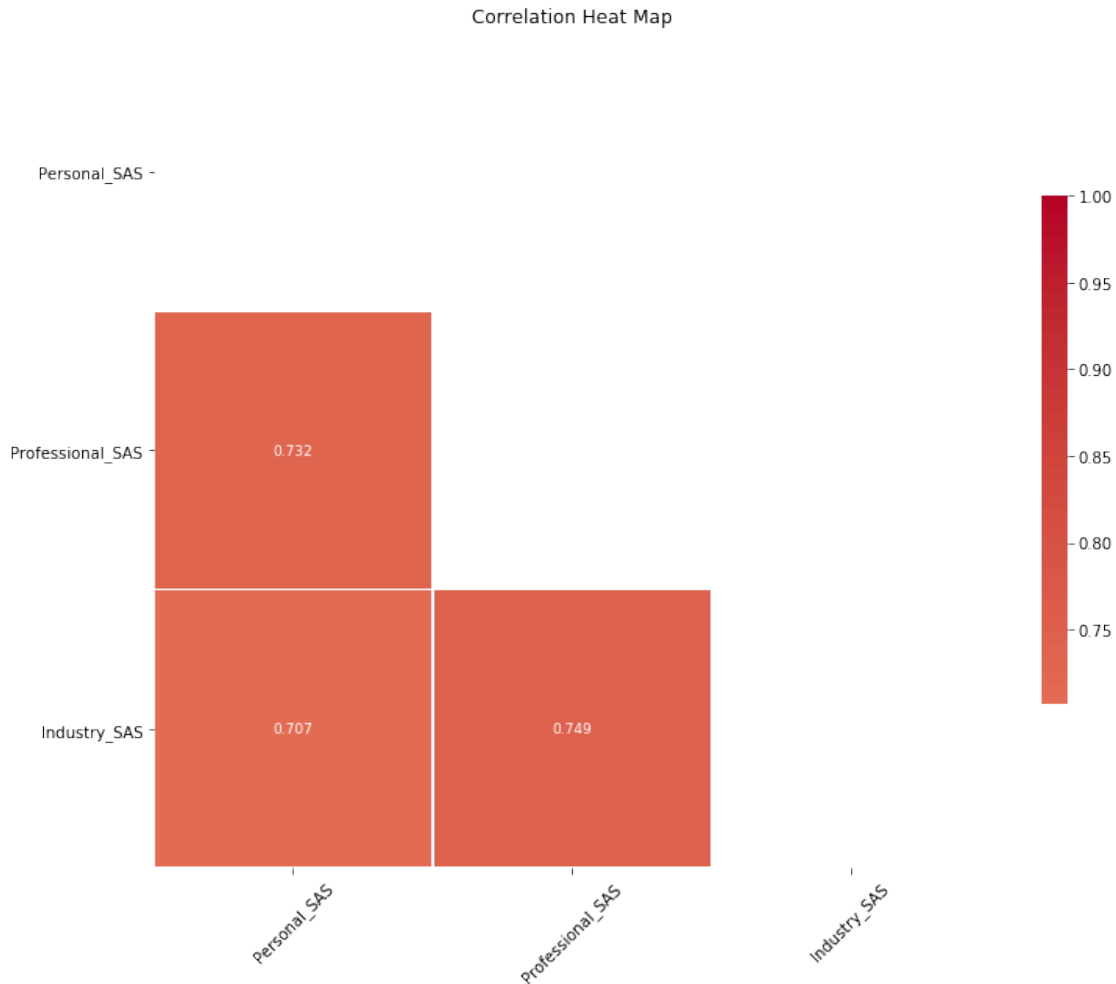


<Figure size 432x288 with 0 Axes>

Correlation Heat Map



<Figure size 432x288 with 0 Axes>



```
In [8]: #visualize preferences by number of programming courses taken. Do students with more e
# single scatter plot example
fig, axis = plt.subplots()
axis.set_xlabel('Professional Preference for R')
axis.set_ylabel('Professional Preference for Python')
group = ()
group = df["Courses_Completed_Bucket"]
plt.title('R and Python Perferences')
scatter_plot = axis.scatter(df[df["Courses_Completed_Bucket"]=="0-3"]['Professional_R'],
df[df["Courses_Completed_Bucket"]=="0-3"]['Professional_Python'],
facecolors = 'none',
edgecolors = 'red',
label = '0-3')
scatter_plot = axis.scatter(df[df["Courses_Completed_Bucket"]=="4-6"]['Professional_R'],
df[df["Courses_Completed_Bucket"]=="4-6"]['Professional_Python'],
facecolors = 'none',
edgecolors = 'blue',
```

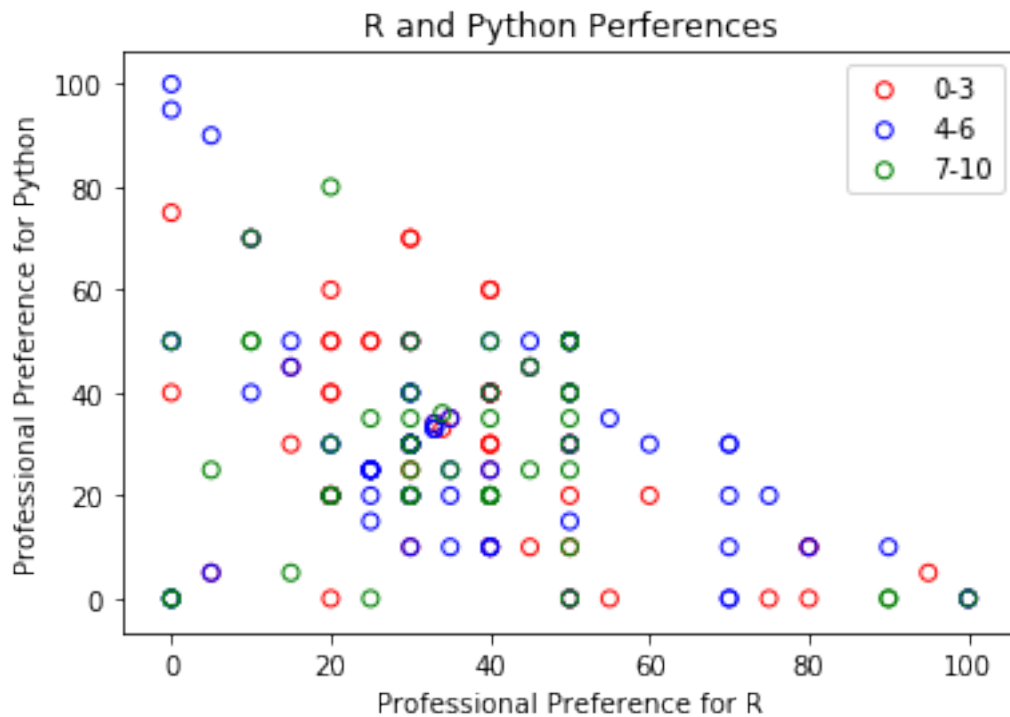
```

label = '4-6')
scatter_plot = axis.scatter(df[df["Courses_Completed_Bucket"]=="7-10"]['Professional_R'],
df[df["Courses_Completed_Bucket"]=="7-10"]['Professional_Python'],
facecolors = 'none',
edgecolors = 'green',
label = '7-10')
plt.legend()

```

*#there doesn't seem to be much a correlation between number of courses taken and preference*

Out [8]: <matplotlib.legend.Legend at 0x18807bf5cf8>



## 0.15 1.4 | Scaling of Variables

```

In [9]: from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import Normalizer
from sklearn.preprocessing import QuantileTransformer

X = df[["Professional_Python"]].dropna()
unscaled_fig, ax = plt.subplots()
sns.distplot(X).set_title('Unscaled')
#unscaled_fig.savefig('Transformation-Unscaled' + '.pdf',

```

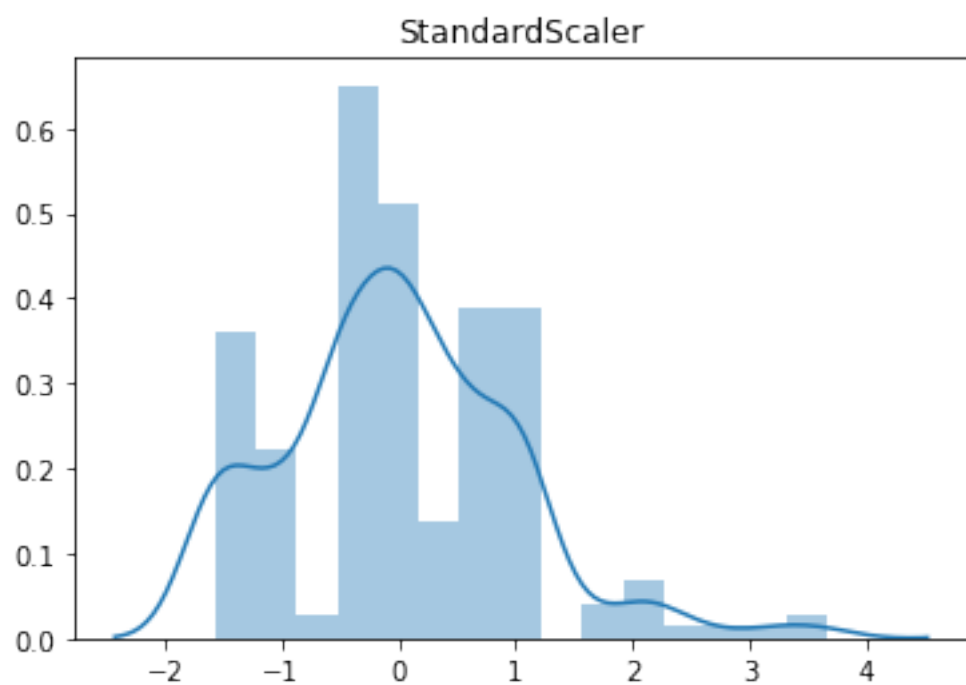
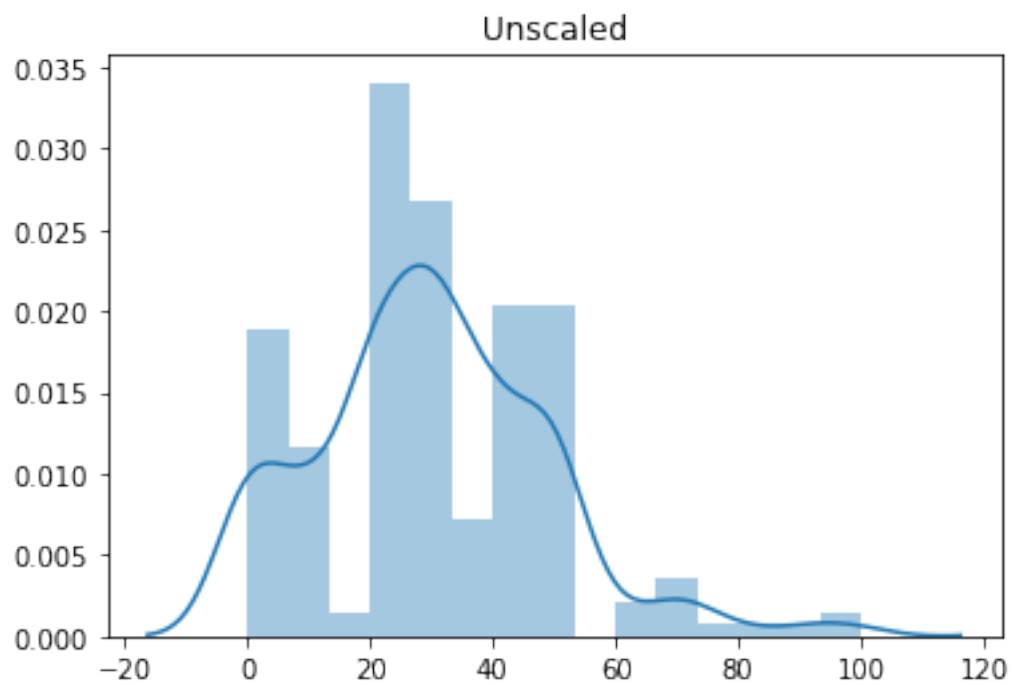
```
#    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
#    orientation='portrait', papertype=None, format=None,
#    transparent=True, pad_inches=0.25, frameon=None)
```

```
standard_fig, ax = plt.subplots()
sns.distplot(StandardScaler().fit_transform(X)).set_title('StandardScaler')
#standard_fig.savefig('Transformation-StandardScaler' + '.pdf',
#    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
#    orientation='portrait', papertype=None, format=None,
#    transparent=True, pad_inches=0.25, frameon=None)
```

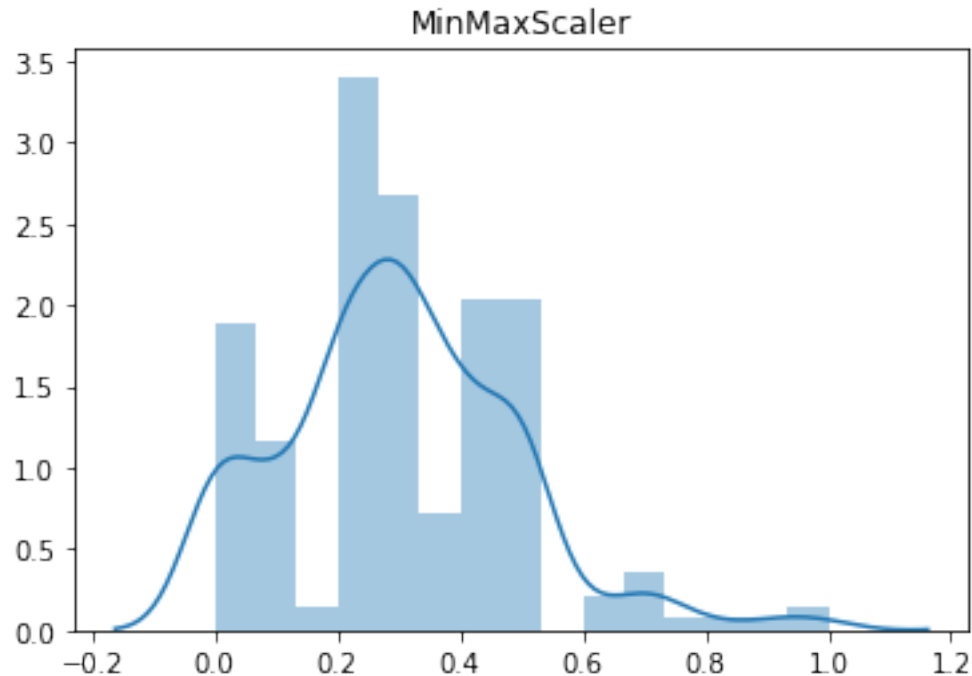
```
minmax_fig, ax = plt.subplots()
sns.distplot(MinMaxScaler().fit_transform(X)).set_title('MinMaxScaler')
minmax_fig.savefig('Transformation-MinMaxScaler' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

*#plotting a variables against a few scaling options. StandardScaler and MinMaxScaler s  
#but StandardScaler is based around 0 which is more intuitive. Additionally, the value  
#extreme outliers, which if they were minmax might be be a better option.*

```
C:\Users\David\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:625: DataConversionWarning:
    return self.partial_fit(X, y)
C:\Users\David\Anaconda3\lib\site-packages\sklearn\base.py:462: DataConversionWarning: Data wi
    return self.fit(X, **fit_params).transform(X)
C:\Users\David\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:323: DataConversionWarning:
    return self.partial_fit(X, y)
```







```
In [16]: #scale all numeric values in data
df2 = pd.DataFrame(StandardScaler().fit_transform(df[personal_fields+professional_fie
columns = df[personal_fields+professional_fields+industry_fields+c
df2.head()

C:\Users\David\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:625: DataConversionWarning:
    return self.partial_fit(X, y)
C:\Users\David\Anaconda3\lib\site-packages\sklearn\base.py:462: DataConversionWarning: Data wi
    return self.fit(X, **fit_params).transform(X)
```

```
Out[16]:
```

	Personal_JavaScalaSpark	Personal_JavaScriptHTMLCSS	\
RespondentID			
5135740122	-0.892507	-0.711586	
5133300037	-0.011911	0.771781	
5132253300	0.868684	-0.711586	
5132096630	-0.011911	0.771781	
5131990362	0.868684	-0.711586	

	Personal_Python	Personal_R	Personal_SAS	\
RespondentID				
5135740122	-2.015302	0.885401	2.454294	
5133300037	1.203583	-0.490044	-1.223949	
5132253300	0.559806	0.197679	-1.223949	
5132096630	-0.405860	-0.146183	0.247349	

5131990362	-2.015302	2.260846	-0.488300
------------	-----------	----------	-----------

	Professional_JavaScalaSpark	Professional_JavaScriptHTMLCSS	\
RespondentID			
5135740122	-0.704282	-0.541476	
5133300037	1.198934	1.776256	
5132253300	-0.704282	-0.541476	
5132096630	0.057005	0.385617	
5131990362	0.818291	-0.541476	

	Professional_Python	Professional_R	Professional_SAS	...	\
RespondentID					
5135740122	-1.572322	-0.548894	3.009440	...	
5133300037	-0.001518	-0.789311	-0.982832	...	
5132253300	0.522084	0.172357	0.081774	...	
5132096630	-0.263318	-0.068060	0.081774	...	
5131990362	-1.572322	2.095694	-0.982832	...	

	PREDICT453	PREDICT454	PREDICT455	PREDICT456	PREDICT457	\
RespondentID						
5135740122	-0.236902	-0.157329	-0.411693	-0.172774	-0.140372	
5133300037	-0.236902	-0.157329	-0.411693	-0.172774	-0.140372	
5132253300	-0.236902	-0.157329	-0.411693	-0.172774	-0.140372	
5132096630	-0.236902	-0.157329	-0.411693	-0.172774	-0.140372	
5131990362	-0.236902	-0.157329	-0.411693	-0.172774	-0.140372	

	Courses_Completed_Prog	Python_Course_Interest	\
RespondentID			
5135740122	-1.972167	-0.790552	
5133300037	0.248935	-1.798517	
5132253300	-0.639506	0.889391	
5132096630	0.248935	0.385408	
5131990362	-0.639506	-0.454563	

	Foundations_DE_Course_Interest	Analytics_App_Course_Interest	\
RespondentID			
5135740122	0.983034	-0.123356	
5133300037	-0.247489	1.021556	
5132253300	0.367773	1.315124	
5132096630	0.060142	1.021556	
5131990362	-1.478012	-0.446280	

	Systems_Analysis_Course_Interest
RespondentID	
5135740122	-0.108502
5133300037	-0.108502
5132253300	0.190402
5132096630	0.847991

5131990362

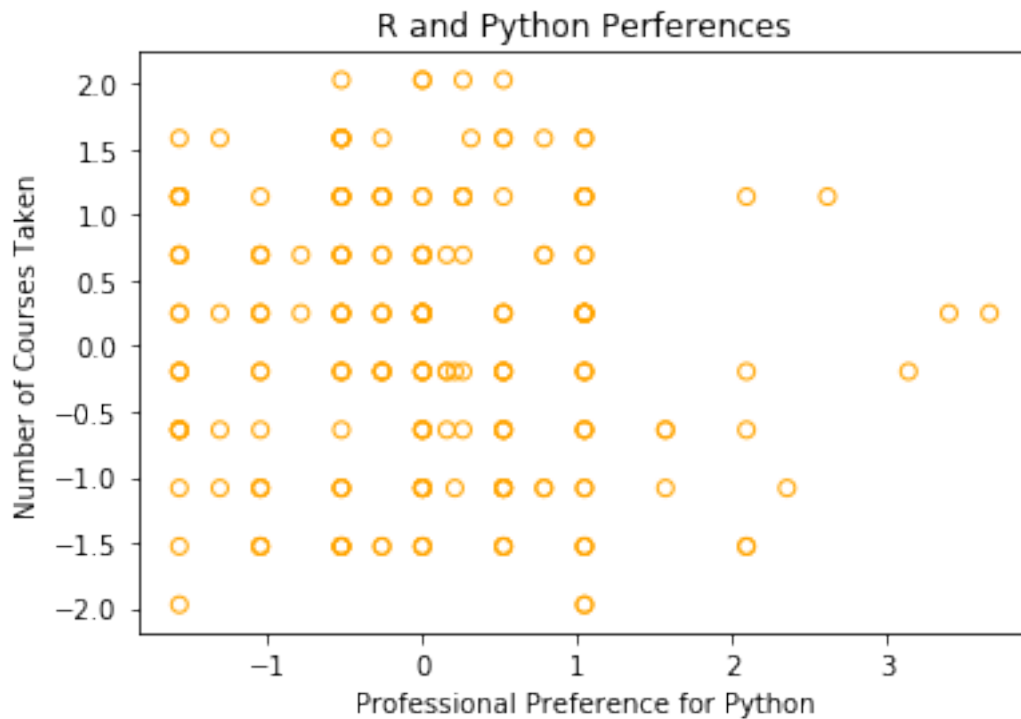
0.788210

[5 rows x 35 columns]

### 0.1.6 1.4.1 | Analysis of Language Preferences (Scaled)

```
In [17]: # with scaled data I can compare two variables that before had two very different ranges
# single scatter plot example
fig, axis = plt.subplots()
axis.set_xlabel('Professional Preference for Python')
axis.set_ylabel('Number of Courses Taken')
plt.title('R and Python Preferences')
scatter_plot = axis.scatter(df2['Professional_Python'],
                             df2["Courses_Completed_Prog"],
                             facecolors = 'none',
                             edgecolors = 'orange')
```

*#this further cements the prognosis that there are no correlations between courses taken and professional preference for Python*



### 0.1.7 1.5 | Analysis of New Courses

```
In [18]: df[interest_fields].head()
```

```
Out [18]:
```

	Python_Course_Interest	Foundations_DE_Course_Interest	\
RespondentID			
5135740122	50.0		90.0
5133300037	20.0		50.0
5132253300	100.0		70.0
5132096630	85.0		60.0
5131990362	60.0		10.0

	Analytics_App_Course_Interest	Systems_Analysis_Course_Interest
RespondentID		
5135740122	51.0	50.0
5133300037	90.0	50.0
5132253300	100.0	60.0
5132096630	90.0	82.0
5131990362	40.0	80.0

```
In [19]: df[interest_fields].describe()
```

```
Out [19]:
```

	Python_Course_Interest	Foundations_DE_Course_Interest	\
count	206.000000	200.000000	
mean	73.529126	58.045000	
std	29.835429	32.588079	
min	0.000000	0.000000	
25%	53.000000	29.500000	
50%	82.500000	60.000000	
75%	100.000000	89.250000	
max	100.000000	100.000000	

	Analytics_App_Course_Interest	Systems_Analysis_Course_Interest
count	203.000000	200.000000
mean	55.201970	53.630000
std	34.147954	33.539493
min	0.000000	0.000000
25%	25.000000	21.500000
50%	60.000000	51.500000
75%	85.000000	80.250000
max	100.000000	100.000000

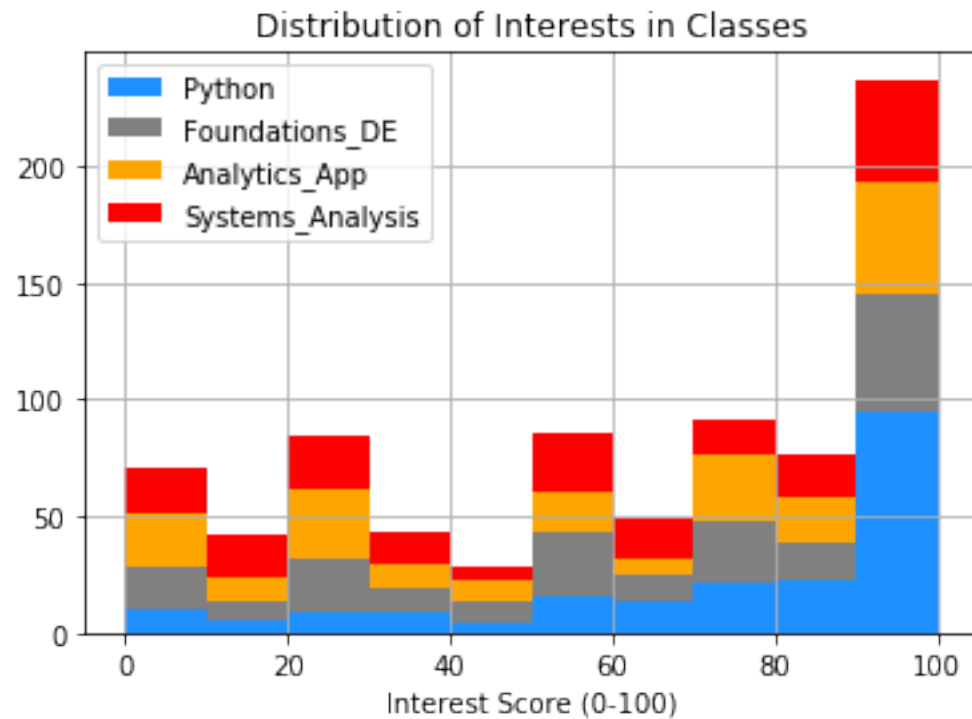
```
In [20]: #isolate age values by class
x1 = df['Python_Course_Interest']
x2 = df['Foundations_DE_Course_Interest']
x3 = df['Analytics_App_Course_Interest']
x4 = df['Systems_Analysis_Course_Interest']

#combine threads and plot
plt.hist([x1,x2,x3,x4],stacked = True,color = ["dodgerblue","grey","orange","red"])
plt.title("Distribution of Interests in Classes")
plt.xlabel("Interest Score (0-100)")
```

```
plt.legend({'Python': 'dodgerblue', 'Foundations_DE': "grey", 'Analytics_App': "orange", 'Systems_Analysis': "red"})
plt.grid(True)
plt.show()
```

*#It looks like the Python course has garnered the most interest receiving the highest scores.*

```
C:\Users\David\Anaconda3\lib\site-packages\numpy\lib\histograms.py:824: RuntimeWarning: invalid value encountered in less
  keep = (tmp_a >= first_edge)
C:\Users\David\Anaconda3\lib\site-packages\numpy\lib\histograms.py:825: RuntimeWarning: invalid value encountered in less
  keep &= (tmp_a <= last_edge)
```



```
In [21]: #isolate age values by class
x1 = df[df["Courses_Completed_Bucket"]=="0-3"]['Python_Course_Interest']
x2 = df[df["Courses_Completed_Bucket"]=="4-6"]['Python_Course_Interest']
x3 = df[df["Courses_Completed_Bucket"]=="7-10"]['Python_Course_Interest']

#combine threads and plot
plt.hist([x1,x2,x3],stacked = True,color = ["dodgerblue","grey","orange"])
plt.title("Distribution of Interests by Courses Taken")
plt.xlabel("Interest Score (0-100)")
plt.legend({'0-3': 'dodgerblue', '4-6': "grey", '7-10': "orange"})
plt.grid(True)
```

```
plt.show()
```

*#It looks like the Python course has garnered the most interest receiving the highest # 100 scores.*

